

Solutions Lecture 6 (Chapter 8)

Make sure to import Numpy, SciPy and Matplotlib to be able to complete all the exercises.

```
import numpy as np
import scipy.optimize as optimize
import matplotlib.pyplot as plt
```

Question 1

In this exercise we will construct the Box-Muller transform that allows the generation of random samples of the normal distribution using samples from the uniform distribution.

If U_0 and U_1 are uniform random variables on $[0, 1]$, then

$$Z_0 = \sqrt{-2 \cdot \ln(U_0)} \cos(2\pi U_1)$$

and

$$Z_1 = \sqrt{-2 \cdot \ln(U_0)} \sin(2\pi U_1)$$

are random variables that follow a normal distribution.

- Write a function `unif_to_norm` that takes as input parameters m and n and returns two $m \times n$ two-dimensional arrays z_0 and z_1 , where `z_0[i, j]` is a sample according to the formula Z_0 and every `z_1[i, j]` a sample according to the formula of Z_1 . That is, you should generate samples from the uniform distribution and transform them to samples from the normal distribution using the formulas above.

```
def unif_to_norm(m,n):
    u_0 = np.random.rand(m,n)
    u_1 = np.random.rand(m,n)

    z_0 = np.sqrt(-2*np.log(u_0))*np.cos(2*np.pi*u_1)
    z_1 = np.sqrt(-2*np.log(u_0))*np.sin(2*np.pi*u_1)
    return z_0, z_1
```

- See if your function works for $m = 10$ and $n = 100$ and do a sanity check by computing the mean (should be close to zero) and standard deviation (should be close to 1) of all elements in `z0` and `z1`.

```
m, n = 10, 100
z0, z1 = unif_to_norm(m,n)
```

```
Mean of elements in z0 is: -0.050313679741913885
Standard deviation of elements in z0 is: 1.029247016960882
```

```
Mean of elements in z1 is: -0.027201045159231253
Standard deviation of elements in z1 is: 0.9887266653281195
```

Question 2

In this exercise, we will write a function that mimics `np.random.choice()` with a specified distribution x in the keyword argument `p`.

- a) Write a function `index_sample(x)` that, for given nonnegative input vector $x = [x_0, \dots, x_{n-1}]$ with $\sum_i x_i = 1$, returns an index $i \in \{0, \dots, n-1\}$ with probability x_i . As a source of randomness, you are only allowed to use the `np.random.rand()` function. Do not use for-loops. Hint: Identify every index with a subinterval of $[0, 1]$, and recall that `np.argmax()` applied to a Boolean vector returns the location of the first `True` element.

```
# Fix randomness
np.random.seed(3)

# Probabilities
x = np.array([1/4, 1/10, 1/10, 1/20, 1/4, 1/4])

# Randomly generated index in {0,1,2,3,4,5} according to x
index = index_sample(x)
print(index)
```

4

```
def index_sample(x):
    # Input: Normalized vector x
    # Output: Index sampled according to probabilities x_i

    u = np.random.rand()
    return np.argmax(u < np.cumsum(x))
```

- b) Vectorize your function to `index_samples(x, k)` so that it can return k sampled indices according to the probabilities in x . Do not use for-loops.

```
# Fix randomness
np.random.seed(3)

# Probabilities
x = np.array([1/4, 1/10, 1/10, 1/20, 1/4, 1/4])
k = 10

index = index_samples(x, k)
```

```

print(index)

[4 4 1 4 5 5 0 0 0 2]

def index_samples(x,k):
    # Input: Normalized vector x, number of samples k
    # Output: k indices sampled according to probabilities x_i

    u = np.random.rand(k,1)
    return np.argmax(u < np.cumsum(x),axis=1)

```

Question 3

The probability density function (pdf) of the exponential distribution $\text{Exp}(\beta)$ with scale parameter $\beta > 0$ is given by

$$f_\beta(x) = \begin{cases} \frac{1}{\beta}e^{-\frac{1}{\beta}x} & x > 0 \\ 0 & x \leq 0 \end{cases}. \quad (1)$$

First, let $X_i \sim \text{Exp}(\beta_i)$ for $i = 1, \dots, n$ and let $X_{\min} = \min_i X_i$.

In this question we will numerically verify the property that the minimum of exponentially random variables is again a random variable with an exponential distribution.

- a) Write a function `min_samples()` that takes as input an array $\beta = [\beta_0, \dots, \beta_{n-1}]$ and a number T . For every β_i it should generate samples x_i^t from X_i for $t = 1, \dots, T$. The function should output an array of length T with the minima

$$\min_{i=0, \dots, n-1} x_i^t$$

for $t = 1, \dots, T$. Do not use for-loops. Hint: Both the `scale` and `size` parameters of `np.random.exponential` allow arrays as input; see the documentation.

```

np.random.seed(3)

beta = np.array([5,2,3])
T = 10

min_data = min_samples(beta,T)
print(min_data)

```

```
[1.03129621 3.57520153 0.15851674 0.0606632  0.65281055 0.04854826
 1.00023611 0.51189432 1.09742183 0.98195207]
```

```

def min_samples(beta,T):
    samples = np.random.exponential(scale=beta,size=(T,np.size(beta)))
    return np.min(samples,axis=1)

```

```

np.random.seed(3)

beta = np.array([5,2,3])
T = 10

min_data = min_samples(beta,T)
print(min_data)

```

[1.03129621 3.57520153 0.15851674 0.0606632 0.65281055 0.04854826
1.00023611 0.51189432 1.09742183 0.98195207]

- b) Write a function `exp_fit()` that takes as input an array of numbers, and fits an exponential distribution (of the form above) to these numbers. Do this by using the `fit()` method for an exponential distribution object in `stats`, which you can read about in the documentation. The output of your function should be the scale parameter of the fitted exponential distribution, where the exponential distribution is as introduced in the beginning of this question. Make sure you lock the location parameter to be 0.

```

np.random.seed(3)

beta = np.array([5,2,3])
T = 10000000
min_data = min_samples(beta,T)

scale = exp_fit(min_data)
print(scale)

```

0.9680113499245512

```

import scipy

def exp_fit(data):
    #expon.fit returns 'loc' and 'scale'. We want loc = 0, hence we choose 'floc=0'
    return scipy.stats.expon.fit(data,floc=0)[1]

beta = np.array([5,2,3])
T = 10000000
min_data = min_samples(beta,T)

scale = exp_fit(min_data)
print(scale)

```

0.9680341501788493

- c) Compute the number

$$\frac{1}{\sum_i \frac{1}{\beta_i}}$$

and observe that it is almost the same as the parameter `scale` in the test output of part b). This means the fitted distribution has a scale parameter β_{\min} satisfying

$$\frac{1}{\beta_{\min}} = \sum_{i=1}^n \frac{1}{\beta_i}.$$

```
beta_min = 1/(np.sum(1/beta))

print(beta_min)
```

0.9677419354838711

Question 4

Suppose you want to sell a laptop to a potential buyer. The valuation the buyer has for the laptop is a nonnegative continuous random variable X with distribution \mathbb{P} . The seller sets a price p for the laptop; the buyer will buy the laptop if the realization of X exceeds the price p .

The revenue of the seller is given by $p \cdot \mathbb{P}(X \geq p)$ if the seller sets price p . The optimal price maximizing the revenue is the so-called *monopoly price* p^* given by

$$p^* = \operatorname{argmax}_p p \cdot \mathbb{P}(X \geq p). \quad (2)$$

The monopoly price p^* can be found (under some assumptions we do not worry about here) by solving the equation

$$\frac{(1 - F(p))}{f(p)} = p \quad (3)$$

where $F(p)$ is the cumulative distribution function (cdf) of X , and $f(p)$ the probability density function (pdf) of X .

- a) Write a function `monopoly_price` that takes as input a continuous distribution (`stats.rv_continuous` object) and a number $\alpha \in (0, 1)$. It should return the monopoly price by solving the equation (3) above. Use the `brentq` method for root finding in your solution with as left-bracket 0, and right-bracket the point \bar{p} for which $F^{-1}(\bar{p}) = \alpha$.

```
# Generate instance of halfnormal distribution
distribution = scipy.stats.halfnorm()

# Set alpha value
alpha = 0.99

# Compute monopoly price
mp = monopoly_price(distribution,alpha)

print(mp)
```

0.7517915246935725

```

def monopoly_equation(p,dist):
    return (1 - dist.cdf(p))/dist.pdf(p) - p

def monopoly_price(dist,alpha):
    a = 0
    b = dist.ppf(alpha)
    result = optimize.root_scalar(monopoly_equation,args=(dist), \
                                    method="brentq",bracket=(a,b))
    return result.root

# Generate instance of halfnormal distribution
distribution = scipy.stats.halfnorm()

# Set alpha value
alpha = 0.99

# Compute monopoly price
mp = monopoly_price(distribution,alpha)

print(mp)

```

0.7517915246935725

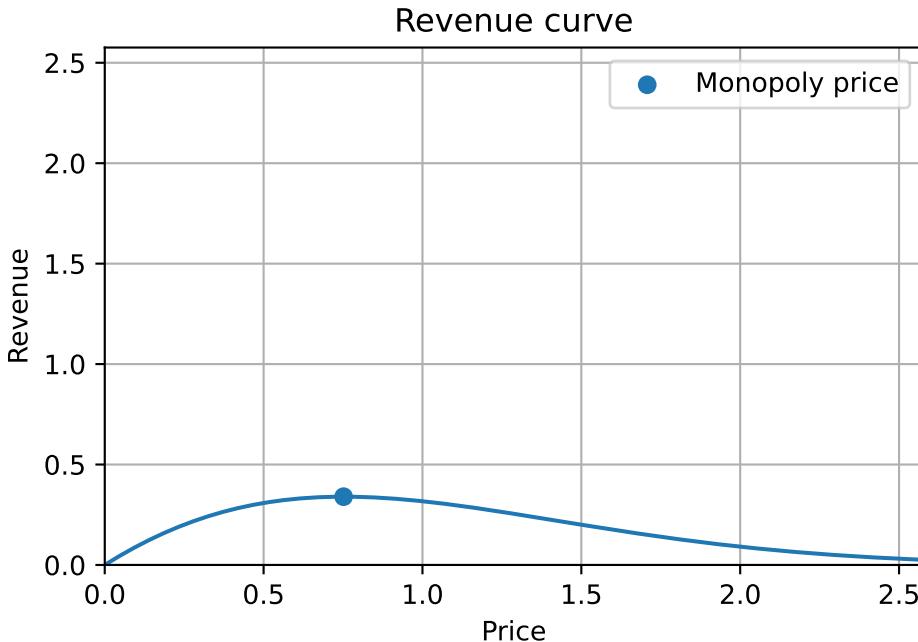
- b) Write a function `revenue_plot` that takes as input a continuous distribution (`stats.rv_continuous` object) and a number $\alpha \in (0, 1)$. It should return a plot of the revenue function $g(p) = p \cdot \mathbb{P}(X \geq p)$ on the interval $[0, \bar{p}]$ with \bar{p} such that $F^{-1}(\bar{p}) = \alpha$. Obtain the monopoly price found using the function `monopoly_price` and add it to your figure with a dot.

```

distribution = scipy.stats.halfnorm()
alpha = 0.99

revenue_plot(distribution,alpha)

```



```

def revenue(p,dist):
    return p*dist.sf(p)

def revenue_plot(dist, alpha):
    b = dist.ppf(alpha) # Set upper bound price range

    p = np.linspace(0,b,1000)

    rev = revenue(p,dist) # Compute revenue curve

    #Create the figure
    plt.figure()

    #Create the plot
    plt.plot(p,rev)

    # Monopoly price
    mp = monopoly_price(dist,alpha)
    plt.scatter(mp,revenue(mp,dist),label='Monopoly price')

    # Create labels for axes
    plt.xlabel('Price')
    plt.ylabel('Revenue')

    # Fix the range of the axes
    plt.xlim(0,b)
    plt.ylim(0,b)

```

```

# Add title to the plot
plt.title('Revenue curve')

# Add legend
plt.legend()

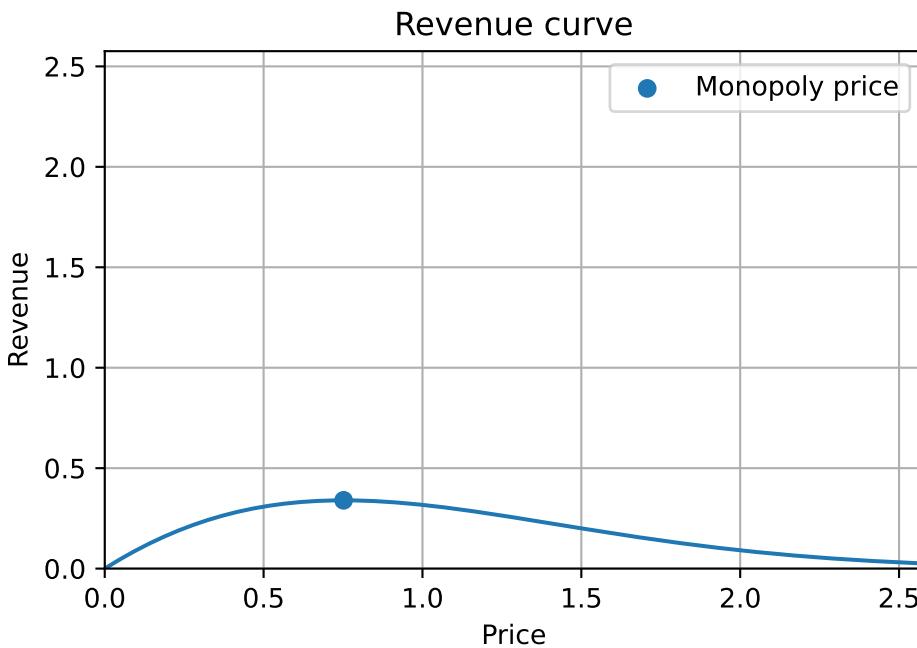
# Add grid to the background
plt.grid()

# Show the plot
plt.show()

distribution = scipy.stats.halfnorm()
alpha = 0.99

revenue_plot(distribution,alpha)

```



Question 5

In this exercise, we will look the problem of sparse vector approximation.¹ You should not use for-loops to answer the questions posed here.

We consider the following setting: We are given an $m \times n$ matrix $A \in [0, 1]^{m \times n}$, a non-negative vector $x = [x_0, \dots, x_{n-1}] \in [0, 1]^n$ with $\sum_i x_i = 1$, and a non-negative vector $y = [y_0, \dots, y_{m-1}] \in [0, 1]^m$ with $\sum_j y_j = 1$.

¹This problem has many applications, for example, in learning theory and in the computation of Nash equilibria in game theory.

We can interpret x and y as discrete probability distributions over the column indices $\{0, \dots, n - 1\}$ and row indices $\{0, \dots, m - 1\}$ of the matrix A , respectively. That is, we have a random variable X that samples column i with prob. x_i for $i = 0, \dots, n - 1$, and a random variable Y that samples row j with prob. y_j for $j = 0, \dots, m - 1$.

The idea of sparse vector approximation is to sample a number of column indices $c_0, \dots, c_{K-1} \in \{0, 1, \dots, n - 1\}$, i.e., K samples of the random variable X , and row indices $r_0, \dots, r_{L-1} \in \{0, 1, \dots, m - 1\}$, i.e., L samples of the random variable Y . Note that the c_i do not have to be distinct. The same holds for the r_j . We then consider for $k = 0, \dots, K - 1$ and $\ell = 0, \dots, L - 1$ the absolute difference

$$|yAx - y^{(\ell)}Ax^{(k)}| \quad (4)$$

where

$$x^{(k)} = \frac{1}{k+1} \sum_{i=0}^k e^{c_i} \quad \text{and} \quad y^{(\ell)} = \frac{1}{\ell+1} \sum_{i=0}^{\ell} (f^{r_i})^T$$

with $e^j \in \{0, 1\}^n$ the unit column vector given by

$$e_i^j = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$

for $j = 0, \dots, n - 1$, and similarly $(f^{r_i})^T \in \{0, 1\}^m$ the row vector with a 1 in position j and zeros elsewhere.

Sparse vector approximation now means that the expression above for $k = K - 1$ and $\ell = L - 1$ converges to zero as K and L grow large, and this happens already for relatively small values (compared to m and n). Informally speaking, $y^{(L)}Ax^{(K)}$ serves as a good approximation to yAx . We will verify this numerically.

- a) Write a function `partial_sums` that for a number n and vector of column indices $c = [c_0, \dots, c_{K-1}] \in \{0, \dots, n - 1\}$, returns a two-dimensional array whose rows are $x^{(0)}, \dots, x^{(K-1)}$. Do not use for-loops.

```
c = [2, 2, 1, 3, 4, 5, 4, 6]
n = 7

x = partial_sums(c,n)
x = np.around(x, decimals=2) # Round to two decimals
print(x)
```

```
[[0.  0.  1.  0.  0.  0.  0.  ]
 [0.  0.  1.  0.  0.  0.  0.  ]
 [0.  0.33 0.67 0.  0.  0.  0.  ]
 [0.  0.25 0.5  0.25 0.  0.  0.  ]
 [0.  0.2  0.4  0.2  0.2  0.  0.  ]
 [0.  0.17 0.33 0.17 0.17 0.17 0.  ]
 [0.  0.14 0.29 0.14 0.29 0.14 0.  ]
 [0.  0.12 0.25 0.12 0.25 0.12 0.12]]
```

```
def partial_sums(c,n):
    K = np.size(c)
```

```

x = np.zeros((K,n))
x[range(K),c] = 1
f = 1/np.arange(1,K+1)[:,None]
x = np.cumsum(x, axis=0)*f
return x

c = [2, 2, 1, 3, 4, 5, 4, 6]
n = 7

x = partial_sums(c,n)
x = np.around(x, decimals=2) # Round to two decimals
print(x)

```

```

[[0.  0.  1.  0.  0.  0.  0.  ]
 [0.  0.  1.  0.  0.  0.  0.  ]
 [0.  0.33 0.67 0.  0.  0.  0.  ]
 [0.  0.25 0.5  0.25 0.  0.  0.  ]
 [0.  0.2  0.4  0.2  0.2  0.  0.  ]
 [0.  0.17 0.33 0.17 0.17 0.17 0.  ]
 [0.  0.14 0.29 0.14 0.29 0.14 0.  ]
 [0.  0.12 0.25 0.12 0.25 0.12 0.12]]

```

- b) Write a function `differences` that takes as input vectors x and y , vectors $y^{(\ell)} \in \mathbb{R}^n$ for $\ell = 0, 1, \dots, L-1$ and $x^{(k)} \in \mathbb{R}^n$ for $k = 0, \dots, K-1$, and an $n \times n$ matrix A , and returns the $K \times L$ matrix $D = (d_{ij})$ whose entries are

$$d_{ij} = |yAx - y^{(\ell)}Ax^{(k)}|$$

Think about an appropriate way to input the vectors yourself. Do not use for-loops in your function.

```

def differences(x,y,X,Y,A):
    # Input: Y is matrix with vectors y^1 on rows,
    #         X is matrix with vectors x^k on rows,
    #         A is square matrix
    # Output: Differences
    return np.abs(y @ (A @ x[:,None]) - Y @ (A @ X.T))

```

- c) Take $n = 100$ and $x = y = \frac{1}{n}(1, 1, \dots, 1)$. Generate resp. $K = 10$ samples of x , corresponding to c_0, \dots, c_9 , and $L = 12$ samples of y , corresponding to r_0, \dots, r_{11} . Take A a randomly generated $n \times n$ matrix with entries in $\{1, 2, 3, \dots, 50\}$ and apply your function from the previous question to it.

```

# Fix randomness
np.random.seed(21)

D = differences(x,y,X,Y,A)
D = np.around(D,decimals=1)

print("D = \n", D)

```

```

D =
[[ 8.5  4.   9.2  6.   6.5  5.2  5.   3.3  1.2  1.8]
 [ 6.5  3.7  0.8  0.2  2.4  3.1  2.3  3.   2.1  3. ]
 [12.5  9.6  7.3  6.1  3.7  1.2  0.6  1.   1.   1.2]
 [ 3.2  4.1  2.1  1.1  0.6  1.8  2.7  3.5  2.6  2.2]
 [ 7.1  6.8  5.1  4.3  2.4  0.4  0.8  2.   1.4  1.4]
 [ 4.6  4.6  4.9  4.4  2.9  1.7  0.7  0.1  0.4  0.1]
 [ 1.9  4.   3.2  3.   2.2  1.5  0.2  0.2  0.1  0.7]
 [ 1.   4.2  3.4  3.8  2.7  1.5  0.1  0.5  0.5  1.3]
 [ 2.3  4.8  3.5  4.1  3.2  2.   0.2  0.3  0.1  1. ]
 [ 0.1  3.3  3.2  3.3  3.   2.1  0.3  0.3  0.1  0.8]
 [ 1.6  2.1  2.6  2.9  2.5  1.8  0.   0.4  0.   0.7]
 [ 3.5  1.7  1.6  1.8  2.   1.1  0.7  1.   0.3  1. ]]
```

```

# Fix randomness
np.random.seed(21)

n = 100

x = np.ones(n)/n
y = np.ones(n)/n

A = np.random.randint(1,51,(n,n))

K = 10
L = 12

c = np.random.choice(range(n),size=K)
r = np.random.choice(range(n),size=L)

X = partial_sums(c,n)
Y = partial_sums(r,n)

D = differences(x,y,X,Y,A)
D = np.around(D,decimals=1)

print("D = \n", D)
```

```

D =
[[ 8.5  4.   9.2  6.   6.5  5.2  5.   3.3  1.2  1.8]
 [ 6.5  3.7  0.8  0.2  2.4  3.1  2.3  3.   2.1  3. ]
 [12.5  9.6  7.3  6.1  3.7  1.2  0.6  1.   1.   1.2]
 [ 3.2  4.1  2.1  1.1  0.6  1.8  2.7  3.5  2.6  2.2]
 [ 7.1  6.8  5.1  4.3  2.4  0.4  0.8  2.   1.4  1.4]
 [ 4.6  4.6  4.9  4.4  2.9  1.7  0.7  0.1  0.4  0.1]
 [ 1.9  4.   3.2  3.   2.2  1.5  0.2  0.2  0.1  0.7]
 [ 1.   4.2  3.4  3.8  2.7  1.5  0.1  0.5  0.5  1.3]]
```

```
[ 2.3  4.8  3.5  4.1  3.2  2.   0.2  0.3  0.1  1. ]
[ 0.1  3.3  3.2  3.3  3.   2.1  0.3  0.3  0.1  0.8]
[ 1.6  2.1  2.6  2.9  2.5  1.8  0.   0.4  0.   0.7]
[ 3.5  1.7  1.6  1.8  2.   1.1  0.7  1.   0.3  1. ]]
```